

Image Classification Using CNN

CS 6301: Practical Aspects of Data Science

Summer 2020: Assignment 2b

Team Member: -
Rajarshi Chattopadhyay (RXC170010)

Introduction

This project involves using Convolutional Neural Network for image classification.

A TensorFlow technique for image classification has been used in with the help of the Keras API of TensorFlow in Python.

Data

CIFAR-10 - Object Recognition in Images <http://www.cs.utoronto.ca/~kriz/cifar.html>
The data is also stored in <https://www.utdallas.edu/~rxc170010/cifar-10-batches-py/>

The CIFAR-10 is labeled subset of the [80 million tiny images](#) dataset, consisting of 60000 32x32 color images in 10 classes, with 6000 images per class.

The data, which comes as vectors, has been reshaped to 32x32x3, indicating a 32x32 image with 3 channels for RGB. This is necessary for using as input for the 2D Convolutional layers of the model.

Training examples 45000, Testing examples 5000

Models

Models were created with different architectures

- The number of neurons in each kernel and the number of hidden layers were modified to test and improve the model.
- Row and column size from input to hidden layer has been kept the same.
- Batch normalization was used in the input layer to ensure there is not much covariance shift from its output.
- Max-pooling was used in the hidden Conv layer to reduce row and column dimensions to half.
- Dropout was used to inactivate randomly 50% of the nodes in Conv layers and 40% in Dense layers for each data example at the beginning. Then changed later.

Library used: Keras https://www.tensorflow.org/guide/keras/sequential_model

Training

The models were trained and tuned to find the best classifier

- Validation was done during training with 33% of the training data
- Batch size of 32 was used
- Number of epochs were changed to improve the model
- Early stopping was done to stop training if there is no significant improvement in validation accuracy after every 3 epochs.
- Checkpointing has been performed to save the best performing model for later use during prediction.

Testing

Each model was tested with test data to obtain the test accuracy and loss.

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_2 (Conv2D)	(None, 30, 30, 16)	4624
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 16)	0
dropout_1 (Dropout)	(None, 15, 15, 16)	0
flatten_1 (Flatten)	(None, 3600)	0
dense_1 (Dense)	(None, 10)	36010

Total params: 41,658

Trainable params: 41,594

Non-trainable params: 64

Train on 30149 samples, validate on 14851 samples

Epoch 1/10

30149/30149 [=====] - 59s 2ms/step - loss: 1.6011 - accuracy: 0.4320 - val_loss: 1.3811 - val_accuracy: 0.5109

Epoch 2/10

30149/30149 [=====] - 63s 2ms/step - loss: 1.2873 - accuracy: 0.5468 - val_loss: 1.2145 - val_accuracy: 0.5792

Epoch 3/10

30149/30149 [=====] - 53s 2ms/step - loss: 1.1631 - accuracy: 0.5951 - val_loss: 1.1732 - val_accuracy: 0.5930

Epoch 4/10

30149/30149 [=====] - 52s 2ms/step - loss: 1.0952 - accuracy: 0.6183 - val_loss: 1.1880 - val_accuracy: 0.5921

Epoch 5/10

30149/30149 [=====] - 51s 2ms/step - loss: 1.0558 - accuracy: 0.6308 - val_loss: 1.0777 - val_accuracy: 0.6209

Epoch 6/10

30149/30149 [=====] - 52s 2ms/step - loss: 1.0227 - accuracy: 0.6413 - val_loss: 1.1034 - val_accuracy: 0.6159


Epoch 7/10

30149/30149 [=====] - 51s 2ms/step - loss: 1.0003 - accuracy: 0.6477 - val_loss: 1.2895 - val_accuracy: 0.5585

Epoch 8/10

30149/30149 [=====] - 56s 2ms/step - loss: 0.9786 - accuracy:
0.6576 - val_loss: 1.1175 - val_accuracy: 0.6072
5000/5000 [=====] - 3s 515us/step

- Train result: train_loss: 0.9786 - train_accuracy: 65.7567
- Validation result: val_loss: 1.1175 - val_accuracy: 60.7165
- Test result: test_loss: 1.0774 - test_accuracy: 62.1800



Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_2 (Conv2D)	(None, 30, 30, 16)	4624
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 16)	0
dropout_1 (Dropout)	(None, 15, 15, 16)	0
conv2d_3 (Conv2D)	(None, 13, 13, 8)	1160
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 8)	0
dropout_2 (Dropout)	(None, 6, 6, 8)	0
flatten_1 (Flatten)	(None, 288)	0
dense_1 (Dense)	(None, 10)	2890

Total params: 9,698

Trainable params: 9,634

Non-trainable params: 64

Train on 30149 samples, validate on 14851 samples

Epoch 1/10

30149/30149 [=====] - 63s 2ms/step - loss: 2.0049 - accuracy: 0.2490 - val_loss: 1.7704 - val_accuracy: 0.3698

Epoch 2/10

30149/30149 [=====] - 58s 2ms/step - loss: 1.6745 - accuracy: 0.3790 - val_loss: 1.6322 - val_accuracy: 0.4062

Epoch 3/10

30149/30149 [=====] - 57s 2ms/step - loss: 1.5544 - accuracy: 0.4308 - val_loss: 1.5371 - val_accuracy: 0.4678

Epoch 4/10

30149/30149 [=====] - 58s 2ms/step - loss: 1.4878 - accuracy: 0.4611 - val_loss: 1.4105 - val_accuracy: 0.5221


Epoch 5/10

30149/30149 [=====] - 60s 2ms/step - loss: 1.4362 - accuracy: 0.4788 - val_loss: 1.3933 - val_accuracy: 0.5116

Epoch 6/10

30149/30149 [=====] - 59s 2ms/step - loss: 1.4063 - accuracy: 0.4938 - val_loss: 1.4191 - val_accuracy: 0.5136
 Epoch 7/10
 30149/30149 [=====] - 61s 2ms/step - loss: 1.3797 - accuracy: 0.4994 - val_loss: 1.3658 - val_accuracy: 0.5371
 Epoch 8/10
 30149/30149 [=====] - 60s 2ms/step - loss: 1.3624 - accuracy: 0.5106 - val_loss: 1.3430 - val_accuracy: 0.5461
 Epoch 9/10
 30149/30149 [=====] - 72s 2ms/step - loss: 1.3458 - accuracy: 0.5155 - val_loss: 1.3414 - val_accuracy: 0.5476
 Epoch 10/10
 30149/30149 [=====] - 66s 2ms/step - loss: 1.3257 - accuracy: 0.5251 - val_loss: 1.3000 - val_accuracy: 0.5591
 5000/5000 [=====] - 2s 461us/step

- Train result: train_loss: 1.3257 - train_accuracy: 52.5092
- Validation result: val_loss: 1.3000 - val_accuracy: 55.9087
- Test result: test_loss: 1.2767 - test_accuracy: 56.8000



Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 64)	1792
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 64)	256
conv2d_2 (Conv2D)	(None, 30, 30, 32)	18464
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout_1 (Dropout)	(None, 15, 15, 32)	0
conv2d_3 (Conv2D)	(None, 13, 13, 8)	2312
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 8)	0
dropout_2 (Dropout)	(None, 6, 6, 8)	0
flatten_1 (Flatten)	(None, 288)	0
dense_1 (Dense)	(None, 10)	2890

Total params: 25,714

Trainable params: 25,586

Non-trainable params: 128

Train on 30149 samples, validate on 14851 samples

Epoch 1/20

30149/30149 [=====] - 93s 3ms/step - loss: 2.0687 - accuracy: 0.2219 - val_loss: 1.8164 - val_accuracy: 0.3779

Epoch 2/20

30149/30149 [=====] - 62s 2ms/step - loss: 2.0563 - accuracy: 0.2337 - val_loss: 1.8974 - val_accuracy: 0.3553

Epoch 3/20

30149/30149 [=====] - 57s 2ms/step - loss: 1.6948 - accuracy: 0.3720 - val_loss: 1.6212 - val_accuracy: 0.4274

Not much accuracy increase, so forcefully interrupted training.

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 64)	1792
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 64)	256
conv2d_2 (Conv2D)	(None, 30, 30, 32)	18464
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout_1 (Dropout)	(None, 15, 15, 32)	0
flatten_1 (Flatten)	(None, 7200)	0
dense_1 (Dense)	(None, 10)	72010

Total params: 92,522

Trainable params: 92,394

Non-trainable params: 128

Train on 30149 samples, validate on 14851 samples

Epoch 1/20

30149/30149 [=====] - 95s 3ms/step - loss: 1.5054 - accuracy: 0.4685 - val_loss: 1.3315 - val_accuracy: 0.5767

Epoch 2/20

30149/30149 [=====] - 92s 3ms/step - loss: 1.1457 - accuracy: 0.6012 - val_loss: 1.1035 - val_accuracy: 0.6129

Epoch 3/20

30149/30149 [=====] - 91s 3ms/step - loss: 1.0429 - accuracy: 0.6390 - val_loss: 1.1997 - val_accuracy: 0.6041

Epoch 4/20

30149/30149 [=====] - 90s 3ms/step - loss: 0.9699 - accuracy: 0.6647 - val_loss: 1.1366 - val_accuracy: 0.6239

Epoch 5/20

30149/30149 [=====] - 103s 3ms/step - loss: 0.9235 - accuracy: 0.6766 - val_loss: 1.3078 - val_accuracy: 0.5639

Epoch 6/20

30149/30149 [=====] - 100s 3ms/step - loss: 0.8923 - accuracy: 0.6901 - val_loss: 1.0221 - val_accuracy: 0.6523


Epoch 7/20

30149/30149 [=====] - 93s 3ms/step - loss: 0.8528 - accuracy: 0.7049 - val_loss: 1.0219 - val_accuracy: 0.6587

Epoch 8/20

30149/30149 [=====] - 108s 4ms/step - loss: 0.8297 - accuracy: 0.7094 - val_loss: 1.0418 - val_accuracy: 0.6496
 Epoch 9/20
 30149/30149 [=====] - 96s 3ms/step - loss: 0.7988 - accuracy: 0.7206 - val_loss: 1.0636 - val_accuracy: 0.6480
 Epoch 10/20
 30149/30149 [=====] - 94s 3ms/step - loss: 0.7851 - accuracy: 0.7252 - val_loss: 1.0772 - val_accuracy: 0.6505
 5000/5000 [=====] - 5s 923us/step

- Train result: train_loss: 0.7851 - train_accuracy: 72.5165
- Validation result: val_loss: 1.0772 - val_accuracy: 65.0529
- Test result: test_loss: 1.0292 - test_accuracy: 67.7200



Model: "sequential_5"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 16)	448
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 16)	64
conv2d_2 (Conv2D)	(None, 30, 30, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout_1 (Dropout)	(None, 15, 15, 32)	0
flatten_1 (Flatten)	(None, 7200)	0
dense_1 (Dense)	(None, 10)	72010

Total params: 77,162

Trainable params: 77,130

Non-trainable params: 32

Train on 30149 samples, validate on 14851 samples

Epoch 1/20

30149/30149 [=====] - 48s 2ms/step - loss: 1.5301 - accuracy: 0.4595 - val_loss: 1.3875 - val_accuracy: 0.5305

Epoch 2/20

30149/30149 [=====] - 44s 1ms/step - loss: 1.1513 - accuracy: 0.5968 - val_loss: 1.1007 - val_accuracy: 0.6228

Epoch 3/20

30149/30149 [=====] - 42s 1ms/step - loss: 1.0318 - accuracy: 0.6403 - val_loss: 1.0817 - val_accuracy: 0.6317

Epoch 4/20

30149/30149 [=====] - 42s 1ms/step - loss: 0.9705 - accuracy: 0.6595 - val_loss: 1.0489 - val_accuracy: 0.6392

Epoch 5/20

30149/30149 [=====] - 52s 2ms/step - loss: 0.9173 - accuracy: 0.6815 - val_loss: 1.0012 - val_accuracy: 0.6578

Epoch 6/20

30149/30149 [=====] - 42s 1ms/step - loss: 0.8869 - accuracy: 0.6908 - val_loss: 1.2114 - val_accuracy: 0.5919

Epoch 7/20

30149/30149 [=====] - 48s 2ms/step - loss: 0.8529 - accuracy: 0.6996 - val_loss: 1.0276 - val_accuracy: 0.6486


Epoch 8/20

30149/30149 [=====] - 48s 2ms/step - loss: 0.8353 - accuracy: 0.7086 - val_loss: 0.9819 - val_accuracy: 0.6639

Epoch 9/20

30149/30149 [=====] - 44s 1ms/step - loss: 0.8156 - accuracy: 0.7134 - val_loss: 1.0106 - val_accuracy: 0.6571
 Epoch 10/20
 30149/30149 [=====] - 45s 1ms/step - loss: 0.7968 - accuracy: 0.7188 - val_loss: 0.9798 - val_accuracy: 0.6668
 Epoch 11/20
 30149/30149 [=====] - 44s 1ms/step - loss: 0.7779 - accuracy: 0.7263 - val_loss: 1.0569 - val_accuracy: 0.6435
 Epoch 12/20
 30149/30149 [=====] - 44s 1ms/step - loss: 0.7662 - accuracy: 0.7285 - val_loss: 0.9988 - val_accuracy: 0.6630
 Epoch 13/20
 30149/30149 [=====] - 47s 2ms/step - loss: 0.7578 - accuracy: 0.7326 - val_loss: 1.0258 - val_accuracy: 0.6567
 5000/5000 [=====] - 2s 375us/step

- Train result: train_loss: 0.7578 - train_accuracy: 73.2628
- Validation result: val_loss: 1.0258 - val_accuracy: 65.6656
- Test result: test_loss: 1.0088 - test_accuracy: 67.2400



Model: "sequential_6"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 16)	448
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 16)	64
conv2d_2 (Conv2D)	(None, 30, 30, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout_1 (Dropout)	(None, 15, 15, 32)	0
conv2d_3 (Conv2D)	(None, 13, 13, 16)	4624
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 16)	0
dropout_2 (Dropout)	(None, 6, 6, 16)	0
flatten_1 (Flatten)	(None, 576)	0
dense_1 (Dense)	(None, 100)	57700
dropout_3 (Dropout)	(None, 100)	0
dense_2 (Dense)	(None, 10)	1010

Total params: 68,486

Trainable params: 68,454

Non-trainable params: 32

Train on 30149 samples, validate on 14851 samples

Epoch 1/20

30149/30149 [=====] - 51s 2ms/step - loss: 1.9638 - accuracy: 0.2634 - val_loss: 1.7368 - val_accuracy: 0.4160

Epoch 2/20

30149/30149 [=====] - 57s 2ms/step - loss: 1.6356 - accuracy: 0.3926 - val_loss: 1.5957 - val_accuracy: 0.4448

Epoch 3/20

30149/30149 [=====] - 48s 2ms/step - loss: 1.5193 - accuracy: 0.4417 - val_loss: 1.4372 - val_accuracy: 0.5078

Epoch 4/20


30149/30149 [=====] - 49s 2ms/step - loss: 1.4334 - accuracy: 0.4775 - val_loss: 1.4105 - val_accuracy: 0.5113

Epoch 5/20

30149/30149 [=====] - 48s 2ms/step - loss: 1.3734 - accuracy:
0.5049 - val_loss: 1.3598 - val_accuracy: 0.5305
Epoch 6/20
30149/30149 [=====] - 57s 2ms/step - loss: 1.3400 - accuracy:
0.5260 - val_loss: 1.3120 - val_accuracy: 0.5513
Epoch 7/20
30149/30149 [=====] - 45s 1ms/step - loss: 1.2897 - accuracy:
0.5397 - val_loss: 1.2218 - val_accuracy: 0.5877
Epoch 8/20
30149/30149 [=====] - 44s 1ms/step - loss: 1.2656 - accuracy:
0.5507 - val_loss: 1.1377 - val_accuracy: 0.6084
Epoch 9/20
30149/30149 [=====] - 47s 2ms/step - loss: 1.2416 - accuracy:
0.5556 - val_loss: 1.1500 - val_accuracy: 0.6062
Epoch 10/20
30149/30149 [=====] - 44s 1ms/step - loss: 1.2273 - accuracy:
0.5633 - val_loss: 1.1943 - val_accuracy: 0.5814
Epoch 11/20
30149/30149 [=====] - 44s 1ms/step - loss: 1.2011 - accuracy:
0.5760 - val_loss: 1.1089 - val_accuracy: 0.6322
Epoch 12/20
30149/30149 [=====] - 45s 2ms/step - loss: 1.1911 - accuracy:
0.5772 - val_loss: 1.0602 - val_accuracy: 0.6311
Epoch 13/20
30149/30149 [=====] - 46s 2ms/step - loss: 1.1738 - accuracy:
0.5817 - val_loss: 1.1360 - val_accuracy: 0.6133
Epoch 14/20
30149/30149 [=====] - 49s 2ms/step - loss: 1.1586 - accuracy:
0.5914 - val_loss: 1.0824 - val_accuracy: 0.6341
Epoch 15/20
30149/30149 [=====] - 47s 2ms/step - loss: 1.1547 - accuracy:
0.5895 - val_loss: 1.0969 - val_accuracy: 0.6269
Epoch 16/20
30149/30149 [=====] - 47s 2ms/step - loss: 1.1397 - accuracy:
0.5948 - val_loss: 1.0932 - val_accuracy: 0.6326
Epoch 17/20
30149/30149 [=====] - 45s 1ms/step - loss: 1.1342 - accuracy:
0.6004 - val_loss: 1.0431 - val_accuracy: 0.6395
Epoch 18/20
30149/30149 [=====] - 45s 1ms/step - loss: 1.1246 - accuracy:
0.6001 - val_loss: 1.0215 - val_accuracy: 0.6441
Epoch 19/20
30149/30149 [=====] - 47s 2ms/step - loss: 1.1217 - accuracy:
0.6040 - val_loss: 0.9914 - val_accuracy: 0.6593
Epoch 20/20

30149/30149 [=====] - 51s 2ms/step - loss: 1.1109 - accuracy: 0.6068 - val_loss: 0.9694 - val_accuracy: 0.6709

- Train result: train_loss: 1.1109 - train_accuracy: 60.6786
- Validation result: val_loss: 0.9694 - val_accuracy: 67.0931
- Test result: test_loss: 0.9389 - test_accuracy: 68.8400



Model: "sequential_7"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 16)	448
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 16)	64
conv2d_2 (Conv2D)	(None, 30, 30, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout_1 (Dropout)	(None, 15, 15, 32)	0
conv2d_3 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_2 (Dropout)	(None, 6, 6, 64)	0
flatten_1 (Flatten)	(None, 2304)	0
dense_1 (Dense)	(None, 100)	230500
dropout_3 (Dropout)	(None, 100)	0
dense_2 (Dense)	(None, 100)	10100
dropout_4 (Dropout)	(None, 100)	0
dense_3 (Dense)	(None, 10)	1010

Total params: 265,258

Trainable params: 265,226

Non-trainable params: 32

Train on 30149 samples, validate on 14851 samples

Epoch 1/20

30149/30149 [=====] - 56s 2ms/step - loss: 1.8070 - accuracy: 0.3323 - val_loss: 1.5862 - val_accuracy: 0.4607

Epoch 2/20

30149/30149 [=====] - 58s 2ms/step - loss: 1.4738 - accuracy: 0.4647 - val_loss: 1.3131 - val_accuracy: 0.5437

Epoch 3/20

30149/30149 [=====] - 51s 2ms/step - loss: 1.3139 - accuracy: 0.5307 - val_loss: 1.3009 - val_accuracy: 0.5441

Epoch 4/20
30149/30149 [=====] - 53s 2ms/step - loss: 1.2165 - accuracy: 0.5743 - val_loss: 1.0960 - val_accuracy: 0.6217

Epoch 5/20
30149/30149 [=====] - 74s 2ms/step - loss: 1.1384 - accuracy: 0.6023 - val_loss: 1.0509 - val_accuracy: 0.6368

Epoch 6/20
30149/30149 [=====] - 56s 2ms/step - loss: 1.0744 - accuracy: 0.6253 - val_loss: 0.9934 - val_accuracy: 0.6701

Epoch 7/20
30149/30149 [=====] - 54s 2ms/step - loss: 1.0289 - accuracy: 0.6451 - val_loss: 0.9802 - val_accuracy: 0.6659

Epoch 8/20
30149/30149 [=====] - 52s 2ms/step - loss: 0.9878 - accuracy: 0.6554 - val_loss: 0.9649 - val_accuracy: 0.6732

Epoch 9/20
30149/30149 [=====] - 53s 2ms/step - loss: 0.9650 - accuracy: 0.6651 - val_loss: 0.9414 - val_accuracy: 0.6890

Epoch 10/20
30149/30149 [=====] - 54s 2ms/step - loss: 0.9234 - accuracy: 0.6818 - val_loss: 0.9386 - val_accuracy: 0.6896

Epoch 11/20
30149/30149 [=====] - 51s 2ms/step - loss: 0.8994 - accuracy: 0.6906 - val_loss: 0.9064 - val_accuracy: 0.6999

Epoch 12/20
30149/30149 [=====] - 59s 2ms/step - loss: 0.8823 - accuracy: 0.6955 - val_loss: 0.9184 - val_accuracy: 0.6875

Epoch 13/20
30149/30149 [=====] - 51s 2ms/step - loss: 0.8614 - accuracy: 0.7035 - val_loss: 0.8729 - val_accuracy: 0.7081

Epoch 14/20
30149/30149 [=====] - 51s 2ms/step - loss: 0.8573 - accuracy: 0.7054 - val_loss: 0.8359 - val_accuracy: 0.7205

Epoch 15/20
30149/30149 [=====] - 52s 2ms/step - loss: 0.8347 - accuracy: 0.7105 - val_loss: 0.8732 - val_accuracy: 0.6933


Epoch 16/20
30149/30149 [=====] - 56s 2ms/step - loss: 0.8189 - accuracy: 0.7169 - val_loss: 0.8671 - val_accuracy: 0.7099

Epoch 17/20
30149/30149 [=====] - 52s 2ms/step - loss: 0.7993 - accuracy: 0.7253 - val_loss: 0.8482 - val_accuracy: 0.7136

5000/5000 [=====] - 2s 400us/step

- Train result: train_loss: 0.7993 - train_accuracy: 72.5331
- Validation result: val_loss: 0.8482 - val_accuracy: 71.3555

○ Test result: test_loss: 0.8135 - test_accuracy: 72.6800



Model: "sequential_8"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_2 (Conv2D)	(None, 30, 30, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 64)	0
dropout_1 (Dropout)	(None, 15, 15, 64)	0
conv2d_3 (Conv2D)	(None, 13, 13, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 128)	0
dropout_2 (Dropout)	(None, 6, 6, 128)	0
flatten_1 (Flatten)	(None, 4608)	0
dense_1 (Dense)	(None, 200)	921800
dropout_3 (Dropout)	(None, 200)	0
dense_2 (Dense)	(None, 100)	20100
dropout_4 (Dropout)	(None, 100)	0
dense_3 (Dense)	(None, 10)	1010

Total params: 1,036,286

Trainable params: 1,036,222

Non-trainable params: 64

Train on 30149 samples, validate on 14851 samples

Epoch 1/20

30149/30149 [=====] - 133s 4ms/step - loss: 1.7147 - accuracy: 0.3689 - val_loss: 1.5906 - val_accuracy: 0.4541

Epoch 2/20



30149/30149 [=====] - 109s 4ms/step - loss: 1.3440 - accuracy: 0.5190 - val_loss: 1.3475 - val_accuracy: 0.5113

Epoch 3/20

30149/30149 [=====] - 103s 3ms/step - loss: 1.1793 - accuracy: 0.5855 - val_loss: 1.0930 - val_accuracy: 0.6214

Epoch 4/20
30149/30149 [=====] - 108s 4ms/step - loss: 1.0616 - accuracy: 0.6323 - val_loss: 1.0425 - val_accuracy: 0.6459
Epoch 5/20
30149/30149 [=====] - 108s 4ms/step - loss: 0.9879 - accuracy: 0.6598 - val_loss: 0.9408 - val_accuracy: 0.6758
Epoch 6/20
30149/30149 [=====] - 113s 4ms/step - loss: 0.9294 - accuracy: 0.6806 - val_loss: 0.8778 - val_accuracy: 0.6996
Epoch 7/20
30149/30149 [=====] - 106s 4ms/step - loss: 0.8744 - accuracy: 0.6991 - val_loss: 0.8763 - val_accuracy: 0.6992
Epoch 8/20
30149/30149 [=====] - 107s 4ms/step - loss: 0.8305 - accuracy: 0.7168 - val_loss: 0.8467 - val_accuracy: 0.7102
Epoch 9/20
30149/30149 [=====] - 110s 4ms/step - loss: 0.7885 - accuracy: 0.7292 - val_loss: 0.8107 - val_accuracy: 0.7213
Epoch 10/20
30149/30149 [=====] - 109s 4ms/step - loss: 0.7570 - accuracy: 0.7387 - val_loss: 0.8970 - val_accuracy: 0.6876
Epoch 11/20
30149/30149 [=====] - 116s 4ms/step - loss: 0.7326 - accuracy: 0.7454 - val_loss: 0.8633 - val_accuracy: 0.7061
Epoch 12/20
30149/30149 [=====] - 112s 4ms/step - loss: 0.6994 - accuracy: 0.7584 - val_loss: 0.7937 - val_accuracy: 0.7267
Epoch 13/20
30149/30149 [=====] - 107s 4ms/step - loss: 0.6755 - accuracy: 0.7677 - val_loss: 0.7709 - val_accuracy: 0.7341
Epoch 14/20
30149/30149 [=====] - 108s 4ms/step - loss: 0.6429 - accuracy: 0.7766 - val_loss: 0.8332 - val_accuracy: 0.7170
Epoch 15/20
30149/30149 [=====] - 113s 4ms/step - loss: 0.6303 - accuracy: 0.7806 - val_loss: 0.8178 - val_accuracy: 0.7205
Epoch 16/20
30149/30149 [=====] - 124s 4ms/step - loss: 0.6166 - accuracy: 0.7898 - val_loss: 0.8003 - val_accuracy: 0.7293
5000/5000 [=====] - 5s 1ms/step

- Train result: train_loss: 0.6166 - train_accuracy: 78.9844
- Validation result: val_loss: 0.8003 - val_accuracy: 72.9311
- Test result: test_loss: 0.7737 - test_accuracy: 73.9600



Iteration	Parameters	Accuracy
1	Input layer = 1 Conv Input layer neurons = 32 Hidden layers = 1 Conv Hidden layers neurons = 16 Hidden layer dropout = 0.5 Activation Function = relu Batch size = 32 Epochs = 10, <i>early stopped at 9</i>	Train = 65.76% Validation = 60.72% Test = 62.18%
2	Input layer = 1 Conv Input layer neurons = 32 Hidden layers = 2 Conv Hidden layers neurons = [16, 8] Hidden layers dropout = [0.5, 0.5] Activation Function = relu Batch size = 32 Epochs = 10	Train = 52.51% Validation = 55.91% Test = 56.80%
3	Input layer = 1 Conv Input layer neurons = 32 Hidden layers = 2 Conv Hidden layers neurons = [32, 8] Hidden layers dropout = [0.5, 0.5] Activation Function = relu Batch size = 50 Epochs = 20, <i>interrupted at 4</i>	<i>Interrupted training in between as accuracy was coming very low in range of 25-35%</i>
4	Input layer = 1 Conv Input layer neurons = 64 Hidden layers = 1 Conv Hidden layers neurons = 32 Hidden layers dropout = 0.5 Activation Function = relu Batch size = 50 Epochs = 20, <i>early stopped at 11</i>	Train = 72.52% Validation = 65.05% Test = 67.72%
5	Input layer = 1 Conv Input layer neurons = 16 Hidden layers = 2 Conv Hidden layers neurons = [32, 16] Hidden layers dropout = [0.5, 0.5] Activation Function = relu Batch size = 50 Epochs = 20, <i>early stopped at 14</i>	Train = 73.26% Validation = 65.67% Test = 67.24%
6	Input layer = 1 Conv Input layer neurons = 16 Hidden layers = 2 Conv, 1 Dense Hidden layers neurons = [32, 16], 100 Hidden layers dropout = [0.5, 0.5], 0.4 Activation Function = relu Batch size = 50 Epochs = 20	Train = 60.67% Validation = 67.09% Test = 68.84%


7	Input layer = 1 Conv Input layer neurons = 16 Hidden layers = 2 Conv, 2 Dense Hidden layers neurons = [32, 64], [100, 100] Hidden layers dropout = [0.25, 0.25], [0.4, 0.3] Activation Function = relu Batch size = 50 Epochs = 20, <i>early stopped at 18</i>	Train = 72.53% Validation = 71.35% Test = 72.68%
8	Input layer = 1 Conv Input layer neurons = 32 Hidden layers = 2 Conv, 2 Dense Hidden layers neurons = [64, 128], [200, 100] Hidden layers dropout = [0.25, 0.25], [0.4, 0.3] Activation Function = relu Batch size = 50 Epochs = 20, <i>early stopped at 17</i>	Train = 78.98% Validation = 72.93% Test = 73.96%

Prediction using best model


Predictions are made using the best model. The training accuracy of the best model is ~79%.

- 4597: Predicted: Deer, Actual: Deer
- 4674: Predicted: Automobile, Actual: Automobile
- 2174: Predicted: Truck, Actual: Truck
- 4978: Predicted: Horse, Actual: Horse
- 4172: Predicted: Cat, Actual: Dog
- 3501: Predicted: Airplane, Actual: Bird


Predicted: Deer, Actual: Deer




Predicted: Automobile, Actual: Automobile




Predicted: Truck, Actual: Truck




Predicted: Horse, Actual: Horse



Predicted: Cat, Actual: Dog



Predicted: Airplane, Actual: Bird



The result is good but could be better with an improved model.

More models could be trained to obtain a better classifier, but infrastructure limitations (e.g. lack of a powerful computer) has limited the CNN model training scope.

Summary and Conclusion

CNN performed slightly better than the normal NN models created using Keras in Python. The NN model built using H2O in R underperformed relative to the one built using Keras in Python.

- Python Keras NN Model accuracy: Training ~75%, Test ~62%
- Python Keras CNN Model accuracy: Training ~79%, Test ~74%

Using H2O in R, the Model training was limited due to infrastructure limitations:

- Rectified Linear Units as the activation function
- Neuron architecture of 3072-128-64-10
- 20 epochs

Using Keras in Python, further training of more deep models could not be done due to infrastructure limitations.

A major advantage of CNN over NN is that we do not have to flatten the input images to 1D as they are capable of working with image data in 2D. This helps in retaining the “spatial” properties of images.

That is the reason, the data that was available as vectors, could be used in NN, but for CNN, a 2D image data could be used.

However, the dataset used for this project was available as 1D vectors.

The model that performed the best had the following architecture:

- Input layer Conv2D that takes in image data of size (32,32,3) i.e. 32x32 RGB image data
- Batch normalization layer following the input layer to ensure there is not much covariance shift from its output.
- 2 Conv2D hidden layers with 64 and 128 neurons and ‘relu’ activation function
- Max-pooling was used in the hidden Conv layers to reduce row and column dimensions.
- Dropout was used to inactivate randomly some of the nodes in Conv layers.
- Flatten layer to flatten the 2D data into 1D Vector for the following dense layers.
- 2 Dense layers with 200 and 100 neurons and ‘relu’ activation function.
- Dropout was used to inactivate randomly some of the nodes in Dense layers.
- Dense output layer with 10 neurons and ‘softmax’ activation function that is used to obtain probabilities of each of the 10 class.

The models were trained and tuned to find the best classifier

- Validation was done during training with 33% of the training data
- Batch size of 50 was used during training
- Number of epochs were set to 20
- Early stopping was done to stop training if there is no significant improvement in validation accuracy after every 3 epochs. The training stopped after 17 epochs, as a result.
- Checkpointing has been performed to save the best performing model for later use during prediction.

- More improved model could have been found but further training of more deep models could not be done due to infrastructure limitations.

CNN is found to be very useful for the image classification of its high accuracy.

The CNN follows a hierarchical (sequential) model which works on building a network, giving a fully connected layer where all the neurons are connected to each other and the output is processed.

On the input image (32×32), a filter (say 3×3) can be used to run along all the pixels (rows, columns) of the image which captures the data and then is passed on to the pooling layer where it performs a mathematical computation and gives out a specific result. Here, the filter actually is run across all the values in the pixel matrix and a dot product of the weights and the pixels is calculated.

In data in the convolution layers, is passed into the Rectified Linear Unit (ReLU) activation function to produce corresponding output, which in turn is fed into the next Convolution layer.

Before the output of the hidden convolution layer is fed to the Dense layers, it can be flattened.

Thus, this full 'convoluted' inter-connected network helps in training the model efficiently, making CNN as the best choice for Image Classification.

References

- Image plotting: https://matplotlib.org/api/_as_gen/matplotlib.pyplot.imshow.html
- Keras: https://www.tensorflow.org/guide/keras/sequential_model